

Cloud Computing and Big Data

Josh Felmeden

November 18, 2021

Contents

1	Overview	4
1.1	Economic Driving Factors	4
1.2	Normal Failure	4
1.3	Blank as a service	4
1.4	Impacts of Cloud Services	5
2	IaaS and AWS	5
2.1	Amazon Simple Storage Service	5
2.2	Amazon Elastic Compute Cloud (EC2)	6
2.2.1	Usage	6
2.3	AWS Simple Queue Service SQS and Architecting for Scale-out	7
2.3.1	Mitch Garnaat's Monster Muck Mashup	7
2.4	AWS simpleDB and AWS Relational Databases (RDB)	8
2.4.1	Amazon SimpleDB	8
2.4.2	Amazon Relation Database Service	9
2.5	Availability Zones	9
3	Virtualisation	9
3.1	The hypervisor	10
3.1.1	Simulated Hardware	10
3.1.2	Xen	11
3.1.3	Areas of Virtualisation	11
3.1.4	KVM	11
3.1.5	Nitro Hypervisor	12
4	Containerisation	12
4.1	Namespaces	13
4.2	Control groups	13
4.3	Container images	13
4.4	COW systems	14
4.5	Performance	14
5	Application Orchestration and Kubernetes	14
5.1	Kubernetes	14
5.1.1	Pod	15
5.1.2	Deployments	15
5.1.3	Services	15
6	Cloud Native Applications	16
6.1	Observability	16
6.2	Metrics	16
7	Serverless Systems	17
7.1	The Four Pillars of Serverless	17
7.2	Invocation	17
7.3	Billing	17
7.4	The Four Stumbling Blocks of Serverless	17

8 Scalable Cloud Architectures	18
8.1 Scaling cube	18
8.2 Load balancers	18
8.3 Decoupling services	19
8.3.1 Message topics	19
8.3.2 Service registries	19
8.4 Automation	19
8.5 Sharding	20
8.6 Decentralised Hashing	20
8.7 Database scaling	20
9 Coordination in Distributed Systems	20
9.1 Replication and Linearisability	21
9.2 CAP theorem	21
9.3 Apache Zookeeper	22
10 DevOps	23
10.1 Infrastructure as Code	24
10.2 Configuration Management	24
11 Google's Core Cloud Technologies	25
11.1 MapReduce	25
11.1.1 MapReduce's Hello World	26
11.1.2 MapReduce Fault Tolerance	26
11.2 Google File System (GFS)	27
12 BigTable	27
12.0.1 SSTables	28
12.0.2 Chubby/Paxos	28
12.1 Spanner	28
12.2 F1	29
13 The Hadoop Ecosystem	29
13.1 Hive	30
13.2 Pig	30
13.3 Cascading	30
14 Cloud Databases: NoSQL	30
14.1 Key-Value DBs	31
14.2 Document Databases	31
14.3 Columnar DBs	31
14.4 Graph DBs	32
15 Queues and Real-time Stream Processing	32
15.1 Real-time stream processing	33
16 Spark	34

1 Overview

1.1 Economic Driving Factors

Cloud computing works by charging someone to use a service for a certain amount of time. If, for example, you let someone use your computer, you might charge them for their usage (if you were a real meanie). How that would be worked out is the *operation expenditure* (Opex) and the *capital expenditure* (Capex, cost of the computer). Combining these, we calculate the *total cost of operation* and calculate the cost per day of the lifespan of the computer (in this example). So, if you used the computer for an hour, you would owe the person a 24th of this daily cost.

Now, this might come down to half a penny, but of course, this doesn't exist any more, so we might charge a whole penny instead. This seems like a marginal profit, but in terms of percentages, this is a 100% increase. The fundamentals of this are how cloud computing generates so much income.

The attraction of this is that the users of the services do not have to pay the Capex, and simply pay opex for the rental of the service.

1.2 Normal Failure

Failures in these systems are to be expected. Say the servers you are buying are guaranteed to have a 99.999% 3-year survival rate ('five-nines reliability'). This is good because there is a very high chance that this remains. Now, if you buy 10 of these servers, the probability that you have all of the servers still working is only 99.99%. Taking this to the extremes, if you buy 500,000 servers (this is standard practise for a lot of the big servers these days), the probability that all of them are still working is a measly 1%. Essentially, failure is something that should be normalised.

Modular data centres are used a lot in big data centres. A unit may be left in the shipping container, and this container is removed or added as a container, meaning that if one module fails, another can just be replaced, meaning that the whole centre doesn't collapse.

In the current climate, the modular centres are considered a whole working unit, and this way of thinking was mobilised by a group of Google engineers.

1.3 Blank as a service

- **SaaS: Software as a service**
 - End user application software that is remotely delivered over the internet.
 - Adobe is an example of this (Adobe Creative Cloud)
 - Used to be bought off the shelf as a CD
- **PaaS: Platform as a Service**
 - Developer application software (middleware) functionality is remotely accessible
 - Might provide a particular combination of OS, web-server, data-base and scripting
 - Popular instance is the free 'LAMP stack' (Linux, Apache, MYSQL and PHP)
 - Used to be dominated by *Google*

- **IaaS: Infrastructure as a Service**

- IT infrastructure almost always virtualised and remotely accessible.
- Virtualisation software allows one physical server to be used by multiple users, each on a virtual machine. If one crashes, the others keep running.
- Used to be dominated by AWS (although this is now a much bigger thing).

While Google and Amazon dominated their respective fields, both of these companies have expanded into the other fields. This is very complicated, but the key thing is that both Google and Amazon offer both PaaS and IaaS.

In around 2015, Amazon created **FaaS: function as a service (aka 'serverless')**:

- No server processes visibly running. Pay only for the time spent executing a function
- Unlike PaaS, scale out without increasing number of servers.
- Amazon Lambda is the best known example, although both Google and Microsoft have answers to this.

1.4 Impacts of Cloud Services

Cloud services have revolutionised business in many ways. It is now possible to do tasks that would previously require access to high performance machines. Instead, it can be sent to a big data centre, and this usage is just charged as rental.

Interoperability is a big issue. **Vendor lock-in** is a concern for a lot of clients, and this simply means that once a client is locked into a vendor, it becomes financially or practically unviable to switch to another supplier. This led to an attempt to develop a sense of unity between cloud companies, where companies created the *Open Cloud Manifesto*. A lot of big companies signed up to this, but a lot of the top dogs didn't sign up for this (unsurprisingly, Google and Amazon). This manifesto seemingly no longer exists as an original. As a consumer, this is bad and worrying because vendor lock in is very possible.

2 IaaS and AWS

2.1 Amazon Simple Storage Service

Amazon S3 is a cloud-based persistent storage. It operates independently from other Amazon services. The simple refers to the features, not that it's simple to use. You can store data in the cloud. You also don't store files, you store *objects*, and these are kept in buckets. Objects have a size limit (5Tb) and a max size on a single upload is 5Gb. All buckets share the same namespace, so no sub-buckets.

It's very easy to use; just use a web GUI that is similar to AWS. It also has a command line interface and has scripting. Default storage can be selected (geographically).

S3 is accessed via API, either by SOAP (xml) or REST (http). Wrappers are available to abstract the

API for programmers.

This is just the storage, so now we will look at the computing of data in the cloud.

2.2 Amazon Elastic Compute Cloud (EC2)

This is a remotely accessible virtual network of virtual servers. Usually, EC2 is run with S3 providing the storage.

A single EC2 virtual server, with the chosen OS etc, is an instance. An instance is instantiated from an Amazon Machine Image (AMI)

- One AMI can be cloned n times to create n instances
- You can build your own by cloning an AMI from your local server
- Or, Amazon have a bunch of prebuilt AMIs that you can choose from

EC2 dynamically assigns a unique IP address to each instance, and this IP can be reassigned, perhaps to someone else. The IP can also be static (also known as an Elastic IP address) at a cost.

EC2 instances run in availability zones (AZs), grouped into regions. AZ is similar to a single data centre, guaranteeing an area has 99.95% uptime.

2.2.1 Usage

Some basic API routes as S3, command-line or a bit of GUI too:

- Amazon's own AWS web console
- Various EC2 plug-ins for browsers
- Third-party cloud management tools

Some AMIs are junk or malware, however, so be careful when selecting this.

There are three types of storage:

- Ephemeral local storage in the instance (dies with instance)
- Persistent cloud (S3)
- SAN-style Elastic Block Storage (EBS)
 - Allows user to create volumes from 1Gb to 1Tb
 - Any number of volumes may be mounted from a single instance

S3 is slow, medium-reliable, but super-durable. Never loses data, so is good for DR backups. Instance storage is simple and cheap, but speed can be really poor. EBS is high on everything, but is complex and costly.

There is some autoscaling based on metrics:

- **Cloudwatch**: automated monitoring of EC2 instances. Reveals many statistics such as CPU

utilisation, disk reads/writes and network traffic. Aggregates and stores monitoring data that can be accessed

- **Auto Scaling:** dynamically adds or removes EC2 instances based on CloudWatch metrics. You define conditions upon which you want to scale up or down your EC2 instances. Auto scaling automatically adds or removes the specified amount of Amazon EC2 instances when it detects that the conditions have been met.
- **Elastic Load Balancing:** automatically distributes incoming application traffic across multiple EC2 instances. Better fault tolerance. Elastic Load Balancing detects unhealthy instances within a pool and automatically reroutes traffic to healthy instances until the unhealthy instances have been restored. Customers can enable ELB within a single AZ or across multiple for consistent application performance.

2.3 AWS Simple Queue Service SQS and Architecting for Scale-out

SQS is reliable, loosely-coupled fault-tolerant storage and delivery of messages. It can be between any clients or computers connected to the internet, and senders and recipients do not have to communicate directly. No requirement that either side be always-available or connected to the internet.

A **message** is up to 256Kb of text-data, sent to SQS and stored until it is delivered. A **queue** serves to group related messages together.

SQS is accessible to clients on any HTTP-enabled platform. Messages are stored redundantly over multiple data-centres truly distributed. Unfortunately, this brings about a few down-sides:

- Message retrievals may be incomplete
- Messages may not be delivered quickly (2-10 seconds)
- Messages may be delivered out of order
- Messages may be redelivered

2.3.1 Mitch Garnaat's Monster Muck Mashup

This is a service that converts AVI videos to mp4 using:

- 'Boto' Python interface to AWS
- S3 to store the video files
- EC2 to do the conversion processing
- SQS for inter-process communication

Uses AWS for *scalability* (scale-out not up, not just buying new resources, adding cheap machines to improve ability.)

The basic steps are:

1. Upload a bunch of video files to a S3 bucket
2. For each file, add a msg to SQS input queue
3. On the EC2 instance, repeat this until input queue is empty:

- a) Read message MI from input queue
- b) Retrieve from S3 the video VI specified in MI
- c) Do the conversion creating VO
- d) Store VO in S3
- e) Write message MO to SQS output queue
- f) Delete MI from input queue

This really illustrates how good AWS is at scaling software. It is good because any number of clients can connect to the bucket. If this bucket is 100% full, it doesn't matter, because additional instances can all talk to the same buckets and queues, so the workload is met.

2.4 AWS simpleDB and AWS Relational Databases (RDB)

2.4.1 Amazon SimpleDB

This software provides:

- Reliable storage of structured textual data
- Languages that allows you to store, modify, query, and retrieve data sets
- Automatic indexing of all stored data

SimpleDB provides 3 main resources:

- **Domains:** Highest-level container for related data *items*: queries only search within one domain
- **Items:** a named collection of *attributes* that represent a data object. Each item has a unique name within the domain; items can be created, modified, or deleted; individual attributes within an item can be manipulated
- **Attributes:** an individual category of information within the item, with a name unique for that item. Item has one or more text string values associated with the name

The downside to this type of storage is that it really does only do one data type (textual). If, for example, you wanted to store pi, you would need to store it as a character string ('3', '.', '1').

There is no Database **schema**, meaning if you mistype something, the database will just accept it as a definition, leading to some unfortunate results. It is not a traditional relational database management system:

- SimpleDB items are stored in a hierarchical structure, not a table
- SimpleDB attribute value max size is 1Kb
- SimpleDB data is all stored as text
- The query language is really basic
- SimpleDB is distributed, so data consistency may suffer due to propagation delays

2.4.2 Amazon Relation Database Service

There is also a relational database system, possibly as a response to Azure. You can set up, operate and scale a full MySQL RDBMS without having to worry about infrastructure provisioning, software maintenance, or common DB management tasks, like backups.

The processing power and storage space can be scaled as needed with a single API call and you can fully initiate fully consistent database snapshots at any time. Can import a dump file to get started. Each DB instance exports a number of metrics to CloudWatch including CPU utilisation, etc.

2.5 Availability Zones

Availability zones are clusters of independent data centres that are up to 20 miles apart. They are interconnected using low latency and enable fault isolation and HA.

Choosing which region to use comes down to a few reasons:

- Data sovereignty and compliance
 - Where are you storing user data?
- Proximity of users to data
- Services and feature availability
- Cost effectiveness
 - Each region has differing costs

High Availability (HA) is the ability to minimise service downtime by using redundant components. It also requires service components in at least two AZs.

Fault tolerance is the same as HA, but also the ability to ensure that no service disruption by using active-active architecture meaning that all components are active all the time. This is of course a lot more costly than just having HA.

IaaS may have HA, but FT unlikely. PaaS will usually have HA, but some services offer FT.

3 Virtualisation

The most fundamental type of cloud computing is IaaS compute, and the most fundamental type of this is the virtual machine. It has unmanaged services, which means you control what they do. You create, save or reuse them. They are networked and connected to storage and have certain security systems.

EC2 instances are a form of virtualisation.

- They mostly run on Xeon processors, but also have other processors available.
- There are lots of different tiers available on AWS that use different purposes
- They run in AMLs.

- After creating a virtual machine, you can start, stop, and terminate the machines. Once the machine is terminated, it will not come back.

In virtualisation, there is some virtual memory that points to addresses in physical memory. It is an abstraction of the storage resources, because the virtual memory seems contiguous, but in physical memory it could be in various locations. The operating system manages this and also has hardware support.

Virtual Memory	Virtual Machine
Abstraction of the RAM memory resources	Abstraction of the storage/process/IO resources.
Mapping of program (virtual) memory addresses to physical addresses	Time slicing VM use of virtual memory addresses/CPU/IO registers in physical addresses.
Operating system manages	Operating system/Hypervisor manages
Hardware support (memory management unit)	Hardware support (e.g. Intel VT)

3.1 The hypervisor

This virtualisation is not a new idea. It has been in use since around 1960, but it has been formalised in 1974, where they defined three important properties:

- *Fidelity*: Program gets the same output whether on VM or hardware
- *Performance*: Performs close to physical computer
- *Safety*: Cannot change or corrupt data on physical computer

This was ensured by analysing the instruction sets and identified two instructions that are the most important: **sensitive** instructions (that can change configurations of resources) and **privileged** instructions. Sensitive instructions need to be caught by the operating system, jumping from user mode to kernel mode. Therefore, all sensitive instructions need to be a subset of privileged instructions. This is all handled by the hypervisor.

Therefore, we have the hypervisor in contact with the server hardware, and then on top of this, we have the VM (or multiple machines) that consist of a guest OS, middleware, and apps. This type of VM is called **bare metal**.

Another type, called **hosted**, has a host OS running on top of the server hardware, and then the hypervisor.

The difference between the two is clear; in type two, the hypervisor can write to the host OS meaning that it can be more lightweight, where as type one can be much faster, at the cost of having to contain stuff that is normally handled by the OS.

3.1.1 Simulated Hardware

Full virtualisation is a complete or almost complete simulation of the underlying guest-machine hardware: virtualised guest OS runs as if it were on a bare machine.

The alternative to this is *paravirtualisation* and is only possible when the source code of the OS is available. The guest OS is edited and recompiled to make system calls into the hypervisor API to execute safe rewrites of sensitive instructions. In this example, the hypervisor doesn't simulate hardware.

EC2 offers both of these virtual machines:

- HVM — hardware virtual machine
 - Virtualised set of hardware
 - Can use OS without modification
 - Intel virtualisation technology
- PV — Paravirtualisation
 - Requires OS to be prepared
 - Doesn't support GPU instances

3.1.2 Xen

Xen is a free, open-source hypervisor developed at University of Cambridge. It has multiple modes:

- Paravirtualisation: guest OS recompiled with modifications
- Hardware assisted virtualisation: Intel x86 and ARM extensions
- Widespread
- Not easily virtualisable (17 instructions that violate the sensitive rules above)

3.1.3 Areas of Virtualisation

There are three main areas of responsibility for virtualisation managers:

- CPU virtualisation
 - Guest has exclusive use of a CPU for a period of time
 - CPU state of the first guest is saved, and the state of the next guest is loaded before control is passed to it
- Memory virtualisation
 - Additional layer of indirection to virtual memory
- I/O virtualisation
 - Hypervisor implements a device model to provide abstractions of the hardware

3.1.4 KVM

Converts Linux into a type-1 hypervisor. Memory manager, process scheduler, I/O stack from Linux. The VM is implemented as a regular Linux process. KVM requires CPU virtualisation extensions to handle instructions.

3.1.5 Nitro Hypervisor

- Based on KVM
- Offloads virtualisation functions to dedicated hardware and software
- Hypervisor mainly provides CPU and memory isolation to EC2 instances
- Nitro cards for VPC networking and EBS storage
 - Can handle NVMe SSD for instance and net storage, transparent encryption
 - Nitro hypervisor not involved in tasks for networking and storage
 - In OS Elastic Network Adapter driver
 - Security groups implemented in the NIC
- Can run bare metal

4 Containerisation

A **docker** is essentially a shipping container system for code. They eliminate the problem of running code on loads of different platforms by simply shipping this container to a computer, and then the application can be run locally.

Containers also bring:

- Reproducability
- Portability
- Flexibility
- Isolation

It makes things like experiments really easy, because the runtime environment is always the same.

Container systems use a number of components:

- Linux containers (LXC)
 - Cgroups and namespaces
- Container runtimes
 - Executables that read the container runtime specification, configure the kernel
- Container images
 - Applications
- Container storage
 - Linux storage systems used to store container images on copy-on-write (COW) file systems
- Container registries
 - Web servers used to store container images
- Container engines
 - Container tools used to pull images from container registries and assemble them on the host before creating the runtime specification and running the container runtime
- Container image builders
 - Tools used to create container images

Docker is the most popular container system:

- A container image format spec
- Tools for building container images
- Tools to manage container images
- Tools to manage instances of containers
- A system to share container images
- Tools to run containers

The **Open Container Initiative** (OCI) created an open governance structure creating industry standards around container formats and runtime. The runtime specification and the image specification were also created as part of this.

On the client computer, you interact with the docker via a CLI. An image is built (or pulled or runned) on the docker daemon. These images are a read only template with instructions for creating a Docker container. You can provide a Dockerfile that will determine how the image is run. The **container** itself is a runnable instance of this image. They are *ephemeral*, meaning that unless specific action is taken, any changes that are not stored to a mounted persistent storage volume are lost.

4.1 Namespaces

Namespaces let you virtualise system resources like the file system or netowrking for each container. Each kind of namespace applies to a specific resource and each namespaec creates barriers between processes:

- **pid namespace**: Responsible for isolating hte process (Process ID)
- **net namespace**: Manages network interfaces
- **ipc namespace**: Manages IPC resources (IPC: InterProcess Communiccation)
- **mnt namespace**: Manages the filesystem mount points (MNT: mount)
- **uts namespace**: Isolates kernel and version identifiers and hostname (UTS: Unix Timesharing System).

4.2 Control groups

Cgroups provide a way to limit access to resources such as CPU and memory that each container can use. Each time a new container is named, a cgroup of the same name appears.

4.3 Container images

They are standard TAR files with a base image and layered with differences. The base image contains:

- Rootfs (container root filesystem)
- JSON file (container configuration)

4.4 COW systems

Instead of overwriting data, data is written somewhere else meaning better recovery. There is built in transactions also. There are snapshots, since we just take the current number of written layers and store a reference to that. If you make a modification to this, we can reconfigure any system with this reference point. Using this means that multiple containers can share the same base images, and just use the reference to the configuration they use.

4.5 Performance

Containers are almost always better than virtual machines. While they both emulate infrastructure and encapsulate the tenant, they differ in key ways:

- VMs provide hardware level virtualisation, while containers provide OS level
- VMs need tens of seconds of provisioning, while containers only require milliseconds
- Virtualisation performance is slower than containers in most dimensions except networking
- VM tenants are fully isolated while containers provide process level isolation to tenants.

5 Application Orchestration and Kubernetes

Orchestration is a number of things:

- Running containers at scale requires management tools
- Manage networking volumes, infrastructure
- Automate
 - Fault tolerance, self-healing
 - Auto scaling on demand
 - DevOps
 - Update/rollback without downtime
- The tools available for us to do this are Mesos, Docker swarm, and Kubernetes

5.1 Kubernetes

The features of Kubernetes include:

- Automatic scheduling of work based on resource usage and constraints
- Self-healing: automatic replacement and rescheduling of failed containers
- Service discovery and load balancing
- Automated rollouts and rollbacks

The most important things to get to grips with are the runtime objects: pods, deployments, and services.

5.1.1 Pod

A pod is a set of one or more containers that act as a unit and are scheduled onto a node together. They share a local network and can share file-system volumes.

5.1.2 Deployments

These describe the *desired state* through declarative updates for pods and ReplicaSets. Essentially, you describe what you want, submit this to Kubernetes, and then it is deployed how you want (without actually telling it exactly what you want). The ReplicaSets balance the number of scheduled and running pods (kills or creates).

The deployment is managed via the spec (what we want), the monitors status (current state), and the template (how).

You are able to create a deployment to rollout a ReplicaSet, declare the new state of the pods, rollback to a previous version, or scale up the deployment to cope with more load.

Deployments use **Yaml files** to describe them (similar to Python, in that it uses indentation to infer scope).

5.1.3 Services

Services define networking to access pods consistently. They also expose pods to the external world. They create groupings of pods that can be referred to by name and have a unique IP address and DNS hostname (these by default cluster scope only). The services ensure the pods that are in use are load balanced, and environment variables containing the IP address of each service in the cluster are injected into all the containers.

Services are defined by a *service config file* (also in Yaml).

There are a few service types:

- Influences networking configuration
- Cluster IP
 - Default
 - Service is discoverable/routable only within the cluster
 - kube-proxy watches API service and updates pod IPTables on change
- NodePort
 - Exposes the service on each Node's IP at a static port (the NodePort)
 - Access the service via <NodeIP>:<NodePort>
 - The simplest way to make your service externally accessible

6 Cloud Native Applications

Cloud Native is container based and elastic. AWS's biggest customer is **Netflix**. They roll out so many instances and microservices, and there is roughly 1 application change per second, meaning new features are coming out very frequently.

Microservices are the opposite of 'monolithic applications' and support organisational agility, such as rolling releases and hot swapping, without any down time. This is achieved through horizontal auto-scaling, design for failure, modular design (containerised etc.). APIs are also used, as well as automating things.

A *monolithic* application is built for one big block. Microservices aim to break this apart, where we have modules separated by API gateways.

The advantages of using this is that the services can be deployed to a subset of your users (called *canary* deployment). This allows testing for small groups of people before rolling them out to everyone. These changes can be tiny; changing the size or colour of a button to see if it changes the click-through rate. Another deployment method is called *blue/green*, where you have an example of your user base, and then you flip a switch and everyone has the new features instantly.

Distributed systems are harder to program, since remote calls are slow and always at risk of failure. There is also the issue of **eventual consistency**. Maintaining strong consistency is really difficult for a distributed system, meaning everyone has to manage this.

6.1 Observability

Once you have a massive cluster, how do you keep track of what's going on inside? Observability is a measure of how well internal states of a system can be inferred from knowledge of its external outputs.

Logging is a very useful method for observation. Local log files are useful at first, but as containers are ephemeral (deleted upon termination), if an instance crashes, you won't be able to retrieve the log data. Additionally, the log could grow potentially unlimited. *Aggregated logs* are useful, and is what is used more commonly. It is a central server that records the log requests. Currently, a *side car* approach is used for logstash, meaning they have access to the same mounted volumes etc, and are an aid to the program.

6.2 Metrics

Metrics are more than logs. This checks the CPU, load average, interrupts, etc. and is able to identify bottlenecks. This is *time sensitive*, so is stored in a time sensitive database. There is a plugin for Kubernetes that runs as a pod.

7 Serverless Systems

Firstly, it's important to know that serverless systems **still use servers!!** The essence of the serverless trend is the *absence* of the server concept during software development. This allows developers to focus more time on things that will *differentiate* their business from other businesses, rather than worry about the inner workings of a server. AWS defines serverless application as one that doesn't require you to provision or manage any servers. Unfortunately, this does lend itself to vendor lock-in.

Serverless systems have lots of possible usages, such as computation, data, messaging, user management and logging.

7.1 The Four Pillars of Serverless

1. No server management
2. Flexible scaling
3. No idle costs
4. High availability

When you think of serverless, the most common flavour is AWS Lambda. The Lambda flavour operates with **two sets of permissions**, one for the invocation, and one for the service (or resource).

7.2 Invocation

A Lambda function is invoked via event or schedule. Minimal Amazon Linux data is run containing the required runtime and function. Function provided with JSON invocation context and parameters. The function executes or times out. Result returned if defined. It is typically invoked in a few ms (also called a warm start) but can take around 1000ms sometimes (cold start). The container is retained for a short time to mitigate this latency.

7.3 Billing

Billing for Lambda systems depends on the amount of memory allocated and duration (for invocation). Therefore, the language used for scripting has a big impact on this. For example, Java would cost a lot due to its heavyweight libraries, while Python or Ruby would cost much less, being very lightweight.

7.4 The Four Stumbling Blocks of Serverless

1. Performance Limitations (cold start etc)
2. Monitoring and Logging (closed source, so difficult to expand existing tools)
3. Vendor lock-in — One of the worst forms of vendor lock in
4. Security and Privacy — All data written is encrypted by Lambda, but in multi-tenant cloud systems, there is a lot of suspicion.

8 Scalable Cloud Architectures

Software architecture is a set of structures needed to reason about a system. It includes software components, the relations between them and properties of both. The architecture may be implicit, meaning that it might not be written down, but every software has an architecture. If you build a significant system, you cannot not have an architecture.

Looking at scaling, we can either scale with pure cores, or use multiple servers (called vertical or horizontal scaling respectively). With **vertical scaling**, there will be no need to change your architecture, while **horizontal scaling** will be cheaper.

8.1 Scaling cube

The scale cube only looks at horizontal scaling. There are three dimensions to this cube:

- X axis: Horizontal duplication — Scale by cloning services and data
 - Each clone can do the work of all the other clones. Work is distributed among clones without bias
 - Inefficient compared to alternatives
 - But, very easy to do
- Y axis: Functional decomposition — Scale by splitting different things. Isolating and making scalable individual responsibility of components
 - Needs to be split in the code base
 - More costly than the x axis.
- Z axis: Data partitioning — Scale by splitting similar things. We partition the domains of incoming requests
 - Data partitioning split relative to the client
 - Improves fault tolerance and cache performance
 - More costly than the other two

8.2 Load balancers

Load balancers are used for:

- Distributing requests
- Managing availability
- Performing health checks
- Session affinity (sticky sessions)
- List of backend servers
- Load balancing policy
- SSL termination
- Alternative: DNS or router based

Load balancers can detect denial of service attacks and shut servers down rather than scaling the service up to keep up with the increased traffic (that could end up costing a hell of a lot of money).

Sticky sessions are really useful. If we use a sticky load balancer, we can always route a client's request to the same server instance, perhaps with an IP hash. This sticky session is really good because say you order things on amazon and have a shopping cart full, we want the next request to go to the server that has just been handling this shopping cart. Cookies are managed by load balancers or the application cookie. It enables session replication via shared session DBs or caches. Sometimes, the session is stored with the client, though this can be risky as the internals of the application can become exposed to the client.

There are a few load balancing algorithms, such as:

- Round robin — Simple even task distribution (ignores the difference in work)
- Weighted round robin
- Least connections

Load balancers allow high availability. If one load balancer has a high load, it can switch to a secondary load balancer. Kubernetes provides stateless LB via a service (round-robin by default). Additionally, they have sticky sessions.

8.3 Decoupling services

8.3.1 Message topics

Consumers can subscribe to different message topics, meaning that producers just send a message to a pub/sub (publish/subscribe), and then only the consumers that are subscribed to a topic would receive this message. This generates a disconnect between the producers and the subscribers but also means more scalability because the producers don't need to worry about each individual consumer.

Message queues are also used, and they are *asynchronous*, meaning it can be queued now and then run later, waiting until a consumer is ready to process. This decouples application logic and is also much more scalable, although it introduces latency into the system.

8.3.2 Service registries

Service registries resolve addresses for names and are based on HA, transactional data stores. For example, Apache.

These decoupling services allow for *event driven* algorithms.

8.4 Automation

You can't manually scale instances in your architecture, so you need to automate this. You'll need an elasticity controller that looks at metrics (CPU, mem, disk...). When exceeded, add more nodes, and when it isn't using much, can decommission some of these nodes.

Another way of judging whether or not we need to scale a system is to look at a job queue and scale depending on this.

8.5 Sharding

Sharding is the act of partitioning and storing a dataset in multiple databases. This might occur when a database gets really full, because replicating this data wouldn't solve the problem, or it might also hold some redundant data. This is particularly useful when having multiple caches. If we have multiple caches and each cache has the same data, we will have a lot of cache *misses*, so we are not getting the most bang for our buck.

8.6 Decentralised Hashing

Regular caches distribute evenly across buckets as a function of n . However, if n changes, all objects need redistribution. The solution to this is *consistent hashing*:

- Each peer received similar amounts of keys
- Little reshuffling on peers entering or leaving
- Map each object to a point on the edge of a circle
- Make a node hold keys for a range of consecutive keys
- If a node is removed, its interval is taken over by a neighbour node. All the remaining caches are unchanged.
- If a new node enters, keys are redistributed from a preceding node. All remaining nodes are unchanged.

8.7 Database scaling

Database scaling is one of the harder things to scale. One technique is **read-replicas**, but this is not good for writing. Another example is **sharding** (looked at above).

9 Coordination in Distributed Systems

Errors in distributed systems can either be a **fault**, which is a single component misbehaving, and a **failure**, which is where the whole system stops. If we have a data center with 10000 hard drives, the probability that one hard drive fails per day is almost 1. Additionally, communication and clocks are unreliable. This leads to a lot of potential errors. We can also get process pauses, such as garbage collection.

9.1 Replication and Linearisability

Once a value has been set for a key, all subsequent read operations return that value until a new value is set. This is a system that conforms to *linearisability*.

9.2 CAP theorem

A good cloud might seek to achieve these three things:

- Consistency
 - All nodes should see the same data at the same time
- Availability
 - Node failures do not prevent survivors from continuing to operate
- Partition tolerance
 - The system continues to operate despite network partitions

However, a bloke called Eric Brewer has a theorem that states that a cloud service cannot simultaneously provide these, and there will always be some kind of trade off.

Consider the following:

- Nodes X and Y suffer a partition
- X wants to move forward with a job but gets no reply from Y
- X must now either
 - wait to hear back from Y (sacrificing availability)
 - Proceed without hearing back (threatening consistency)
 - or never be partitioned in the first place (losing partition tolerance)

The CAP theorem was understood to mean that the designer's job is to pick 2 out of 3. But, with cloud, we need partition-tolerance, so you have to choose between C and A. However, these three are not binary concepts. Each is a matter of degree, so we only rule out the possibility of 'perfect' consistency and availability.

Let's now observe some trade-offs:

- **Strong consistency** — After an update completes, any subsequent access will return the same updated value.
- **Weak consistency** — Subsequent accesses are not guaranteed to return the updated value
- **Eventual consistency** — Eventually, all subsequent accesses will return the updated value (e.g. updates will eventually propagate to replicas).

This eventual consistency is good for something like a blog post. But in a different context, it might not be enough.

Prioritizing C could be used where a cloud supports operations in the real world that are really hard to roll back.

We might have a lot of replicas. Now, there are some problems we need to overcome when dealing

with a lot of replicas, and these might be something like where multiple requests come in and it is unsure which one should be processed first. The way around this is to assign each 'problem' a value, and then to choose a valid value at random. Not one that is 'best' or came first, just a valid value. These replicas will have three main roles:

- Proposers: learn values that are already accepted and propose these values
- Acceptors: let proposers know already accepted values, accept or reject proposals, and reach a consensus on choosing a particular proposal or value
- Learners: become aware of the chosen proposal or value and action it

Replicas may play more than one of these roles at different times. Often, a replica is elected to be a 'distinguished' (or privileged) learner/proposer. This is where they are the only one allowed to play the role. Also, different proposals must never use the same ID number (obviously).

Here is how the algorithm unfolds:

- Proposer P picks an ID number N, which is higher than any chosen before.
- P asks some majority of acceptors to prepare for proposal N.
- Only a majority and not all are asked because if a majority of replicas agree on a value, then a different majority cannot agree on a different one.
- At this stage, P hasn't proposed a value, and is only trying to engage enough acceptors
- Acceptors may or may not sign up for this
- If an acceptor A gets a prepare message from P with ID, N, bigger than any seen so far, then A promises to ignore future proposals with ID less than N.
- If A has already accepted a proposal, A tells P the value of the last accepted proposal.
- If N is too low, A ignore P.
- If P gets the majority of promises back, P can set a value.
- If the acceptors have told their accepted values, P chooses the value, v, associated with the highest proposal ID number. If no accepts so far, P sets value.
- P request acceptors accept this proposal
- They must accept unless P was too slow and already proposed another proposer.
- An acceptor will accept the first proposal it sees. System will not hang waiting to compare multiple ideas.
- An acceptor will accept the proposal with the highest ID number. This keeps the system going when proposers die mid proposal.
- A majority of acceptors must accept the same value keeping the system consistent
- Once a value is chosen, all proposals with a higher ID number choose to recommend this same value, meaning we can build a consensus without newcomers disrupting it.

9.3 Apache Zookeeper

Zookeeper is part of a family of linearisable, durable data stores, with a rich client that prioritises consistency over availability. There are typically five nodes in a cluster and are optimised for read, rather than write.

10 DevOps

Software development lifecycle:

1. Planning
2. requirements Definition
3. Software design
4. Software development
5. Software testing
6. Software deployment

Before we had Agile and DevOps, we used to use the waterfall system. The main problem with this was that because it is very one way, it becomes very hard to change things later on in the development. As a result, the final product may not reflect customer requirements and becomes very expensive to change features and spec late in the process.

Agile became the favoured development method in the 1990s. It has iterative development cycles using 'sprints'. Each of these sprints delivers a minimum viable product, each of the sprints usually being 2 weeks long. The product owner is the customer specifying feature requirements for each sprint and signing off releases. Agile methodology allows for quick iteration on feedback.

DevOps breaks the wall down between development and operations. The goal is simple: releasing updates as quickly as possible without losing quality. It's a solution to cultural norms: developers want *agility* while operations want *stability*. DevOps is an extension of the Agile methodology:

- People: Break down Dev and IT Ops silos through improved communications and joint accountability for product quality
- Processes: Increase throughput and quality by automating manual processes
- Tools: Create an end-to-end scripted tool chain (pipeline) for speed and consistency.

DevOps uses cross functional Dev/Ops teams. This allows for:

- Collaboration
 - Sprint dev teams may include QA and IT Ops members
- Responsibility
 - Teams are responsible to end-to-end quality of product
- Awareness
 - Devs learn about potential Ops issues, Ops learn about code
- Communication
 - Teams use integrated comms tools like slack and JIRA to improve issue reporting and response

Another benefit of DevOps is that there are frequent *commits* and *builds*. These builds should be automated so that they can be performed multiple times per day. These processes should also be automated:

- Continuous Integration:
 - Shared source code repository
 - Automated unit testing on commit

- Frequent integration builds and tests
- **Ensures robust code base**
- Continuous delivery
 - CI + more testing + packaging
 - **Ensures production-ready release**
- Optionally, could go further with Continuous deployment
 - CI + CD + automated deployment to prod environment

The third benefit of DevOps is **robust product releases**:

- Fail-Fast philosophy
 - CI/CD automated tests pinpoint code issues early
 - Dev teams alerted immediately
- Integrated monitoring
 - Developers include instrumentation code in app to create metrics data
 - Metrics dashboards reviewed by team for emerging troublespots
- Integrated Security (SecOps)
 - Include automated tests on code vulnerabilities, package integrity, valid certificates in CI/CD

And the fourth and final benefit is **consistent deployment environments**

- Identical Environments
 - Dev/Stage/Prod environments are identical in specification
 - Avoids 'works on my machine!'
 - Use identical machine configs
- Infrastructure as Code (IaC)
 - Templating
 - Config Management
 - Scripting

10.1 Infrastructure as Code

The idea behind IaC is to automate the creation, deletion and updates to infrastructure. It also provides consistent repeatability for multiple environments. It enables secure one-use disposable environments. Version controlling infrastructure is also very easy with this. Performing infra-audits is also easy. All of these things are dependent on robust IaaS APIs.

10.2 Configuration Management

The use of config management greatly simplifies configuration of multiple servers (or nodes).

- Config can be OS patches, apps, websites, scripts, properties ...
- Define config state for a node type
- CM server/client regularly ensures node state is valid and up to date

The config is *idempotent*:

- Resources are only allocated once
- Manual changes are detected and rolled back

Popular CM systems include:

- Chef, Puppet, Ansible, SaltStack, DSC

IaC often uses *templating*, which is infrastructure defined in either a JSON or YAML template file. The template engine parses file and calls IaaS APIs to build a collection of resources (a 'stack'). The most popular IaC templating engines include:

- CloudFormation (AWS)
 - Has a template section (objects to be created) and several optional sections (including parameters, conditions, mappings, etc.)
- Terraform (HashiCorp)

11 Google's Core Cloud Technologies

We are going to look at the 'under the hood' of Google's underlying infrastructure:

- MapReduce (Hadoop)
- Google File System (GFS) (HDFS)
- Big table (Hbase)
- Spanner
- F1

All of these Google technologies are *proprietary*. The equivalents in brackets are the free equivalent, offered by **Apache**. The Hadoop software system does not really have any direct equivalent for Spanner or F1.

11.1 MapReduce

MapReduce is the key that allowed Google to cluster webpages, and is really the secret behind Google's success. MapReduce is a style of programming and of implementation for generating and processing very large data sets. It's not a specific algorithm, but rather more of an architecture. It was created at Google, and since then around 10,000+ programs have been implemented via MapReduce at Google alone.

It was intended to automatically parallelize 'big' data analysis tasks over large clusters built from networks of cheap commodities with some fault-tolerances.

MapReduce partitions the input data into a number of 'splits' (the parallelization we mentioned earlier). The size of these splits are defined by the user on creation of the job. There is a *master* in the interim that assigns work to individual machines (called workers). The **map** part of MapReduce

is where some workers are given responsibility for the input files, while the **reduce** part is where the workers are in charge of the output files.

The master tells the workers on the map side to partition their received data (one of the splits), and this data is going to be read by particular reduce workers.

1. First, there is a read operation, where one of the workers *reads* from one of the splits. This is where they extract key/value pairs out of the data
2. This data is then passed to a *map* function. The map function has been specified by the user, and this creates intermediate k/v pairs; initially stored in a buffer in the worker's memory, before being written to some point in memory.
3. At this point, the worker feeds back to the master the location of the k/v pair in memory.
4. The master forwards this data to the reduce workers
5. This data, among others worked out by other workers, will be read by a reduce worker.
6. This reduce worker iterates over this new data, and passes each unique key and intermediate values to the user's reduce function (again, provided by the programmer).
7. The results of this are written to the relevant output file.

This is, of course, repeated for a lot of different workers and for a lot of different output files.

11.1.1 MapReduce's Hello World

The equivalent of hello world for MapReduce is counting the number of occurrences of each word in a body of text.

The function might look something like:

```
map(String key, String value):  
    //key: document name  
    //value: document contents  
    for each word w in value:  
        EmitIntermediate(w, "1");  
  
Reduce(String key, Iterator values):  
    //key: a word  
    //values: a list of counts  
    int result = 0;  
    for each v in values:  
        results += ParseInt(v);  
    Emit(AsString(result));
```

11.1.2 MapReduce Fault Tolerance

In reality, the instances of MapReduce will have hundreds or thousands of worker machines, so normal failure is a significant issue. To resolve this, a master pings every worker occasionally. If a worker does not respond before a timeout, then the master marks the user as being failed.

Because the results of any map tasks completed by the failed worker, the master resets the failed

server's list of map tasks as unallocated and reschedules them onto other map worker machines, replacing the failed machine.

All reduce workers that still need data from the failed map worker are notified of the replacement map worker(s) and so the reducers switch to reading from the local disks of those replacement map workers.

When a map task finishes, it sends a notification to the Master and includes names of the local temporary files where the map outputs are: if the master receives such a notification for a task that has already been completed by a different worker, it ignores it, else it records the completion data.

11.2 Google File System (GFS)

When Google started, it committed to building a huge scalable distributed storage capability using cheap commodity components. The components used were cheap and therefore unreliable, so 'normal failure' had to be something dealt with by the file storage system. GFS delivers this.

Files in GFS are divided into fixed-sized 64Mb blocks called **chunks**. Each chunk is assigned a unique i.d. called a **chunk handle**. Chunks are stored on **chunk servers**. Fault tolerance is provided by replicating each chunk across multiple servers. Meta-data needs to be recorded so GFS knows what is where.

GFS consists of three main component types:

- GFS Master Server (x1)
- GFS Chunk server (multiple)
- GFS Client (multiple)

12 BigTable

BigTable manages large-scale structured data; designed to reliably scale to petabytes of data spread across thousands of machines. BigTable is widely applicable within Google. Workload types vary from batch processing jobs where data-throughput rates are key. It is also the seminal "NoSQL" database: ultra large scale, but not fully relational. The locality of the data is under the control of the client.

A BigTable *cluster* is a set of processes running BigTable. Each cluster serves a set of *tables*. Each table is a sparse, distributed, persistent, sorted *map*. The map is from three dimensions (row, column, time) onto a string value (a *cell*). The rows contain data sorted in alpha-order by row key. Rows with consecutive keys are grouped into *tablets*. A table can have any number of columns; the keys of which are grouped into sets called *column families*. Different cells can contain multiple versions of the same data differentiated by a *timestamp*. Garbage collection can keep the most recent n versions of data and/or only keep data that is within n seconds of 'now'.

The API of BigTable gives access to functions for:

- put/get of table entries
- creating/deleting tables and column families
- changing metadata associated with clusters, tables, and column families

Clients can provide scripts that are executed on BigTable servers: scripting languages called *Sawzall*; client scripts are unable to write back into BigTable, but they can produce filtered, summarized and transformed data from the table. BigTable can be an input source to, or an output target from, MapReduce. The read/write access to disk is via GFS.

12.0.1 SSTables

BigTable internally uses **sorted-string tables** (SSTable), which are an immutable file format, to store data. There isn't much public information on SSTable, but there is an open-source release of this called **LevelDB**, written in Ruby, that is very similar to how BigTable uses SSTables.

12.0.2 Chubby/Paxos

Chubby is a distributed lock service for shared-resource access synchronisation: a system that prioritises availability and reliability rather than high performance. It implements **Paxos**.

12.1 Spanner

Spanner is Google's globally distributed database. It is so called as it 'spans' the entire globe; data physically located in any data center anywhere on the planet should be seamlessly accessible and *consistent*. NoSQL approaches like BigTable can store data across multiple datacentres, but ensuring the data is consistent in all locations is very hard. It is an issue due to latency between regions.

Timestamping is used to combat this and is usually done via Network Time Protocol (NTP). NTP can sometimes cause problems, however, such as the subtraction of a leap-second that caused several web services to go down.

For Spanner, Google ignored NTP and developed *TrueTime*. Spanner data-centers are equipped with their own atomic clocks, and with GPS receivers to get time reference from GPS signals. The clocks and GPS Rx feed time data to a number of master servers. The masters disseminate time data to all other server machines in the DC. Each server runs a daemon that constantly checks with the masters in the local DC and in other DCs to get the time reference signals and to compute this from a consensus view on what the real time is. This means that data can be committed to DCs at multiple different locations and the common clock means that there is no dispute about which one happened first.

Using Spanner, Google can accurately replicate data across multiple data centres. This gives greater system resilience as any outage can be quickly addressed by firing up the most recent replica or mirror image. It also reduces latency, since if one replica is getting a lot of hits, and a backlog is building up, requests can be diverted to another replica. Spanner automatically shards data across

data-centers around the world and automatically reshards data as the amount of data or number of servers changes. Load balancing is done automatically.

The upfront costs for these clocks and GPS technologies are in the thousands of dollars, but spread across the data centers on a per server basis, they are actually quite cheap.

Spanner offers a number of novel features:

- Applications given fine-grain dynamic control of data replication configs.
- Externally consistent reads and writes
 - External consistency guarantees that a transaction will always receive current information. A system is said to provide external consistency if it guarantees that the schedule it will use to process a set of transactions is equivalent to its external schedule.
- Globally consistent reads across the database at a timestamp.

These features enable Spanner to support a really high scaling system, with things such as consistent backups in bulk.

12.2 F1

F1 is Google's NoSQL database. It is built on Spanner and uses five replicas spread across the US. It is a rewrite of Google's advertising backend.

F1 chose to use Spanner over other solutions for several reasons:

- Spanner removes the need to manually reshard
- Spanner provides synchronous replication and automatic failover
- F1 requires strong transactional semantics and this is not practical in other NoSQL systems.

F1 has several design goals:

- **Scalability:** scale up, trivially and transparently, just by adding resources
- **Availability:** system must never go down
- **Consistency:** provide ACID transactions
- **Usability:** full SQL support, without using SQL

13 The Hadoop Ecosystem

Originally, there was Hadoop, which is Apache's answer to MapReduce. Of course, Hadoop needed something like GFS, which Apache dubbed HDFS. BigTable was needed, and this was called HBase. These three projects were known as 'The Hadoop Core'.

Writing MR jobs is pretty tricky, so some programmatic abstraction was required. This was called Apache Pig. Additionally, they needed an SQL-like query language, which they called Apache Hive.

Mahout is Apache's scalable machine learning platform, and runs on Hadoop/Spark. It provides tools

for using Hadoop to create systems that perform clustering, classification, and collaborative filtering at high speed on large data sets.

Giraph is a graph database system used by Hadoop.

13.1 Hive

Hive is best suited to *data warehousing*. It operates on relatively static data where the response time is not critical. It has its own SQL-like query language called HiveQL/HQL, which is much more concise than lower level MR code.

Routes into Hive could either be a CLI, a web interface, or programmatically. Hive has a **driver** that compiles input, optimises computation and executes subjobs.

13.2 Pig

Pig is actually short for Pig Latin, and works on HDFS and NoSQL. Users describe how data is read/processed/stored in parallel. Dataflows could be linear or DAG (directed, acyclic graph). Interestingly, it has no if/for/while. Unlike SQL/HQL, Pig allows users to describe exactly how data will be processed.

13.3 Cascading

Similar to Hive and Pig, it is another layer on Hadoop, extending its use. It is still in development. Cascading lets you specify plans for queries. It provides a set of operators and tools for combining them into a DAG (as used in Pig). Each vertex in the DAG is an operator, and data flows from operator to operator. Therefore, it's very similar to Pig, except the graph can be displayed.

Once one job is finished, the next queued job is woken up, and then begins executing.

14 Cloud Databases: NoSQL

NoSQL is a really good answer to traditional RDBMs since as databases got bigger, people designed bigger single servers, where costs increase nonlinearly. Even when scaled up, they suffer from 'inelasticity', in that it's hard to scale them back down again when demands reduced. Instead, on clusters of cheap commodity servers and don't need ACID. We'll look at different types of NoSQL databases

14.1 Key-Value DBs

This is the simplest of NoSQL DB types. There are virtually no constraints (schemaless). A consequence of this is that they are often very fast, and don't require a heavy up-front investment in designing schema and are essentially similar to an associative array that allows any value (no typing restrictions).

Keys must be unique within any one of the KV namespace, which may be the whole DB or one of several buckets within a DB.

Values can be pretty much any digital object.

KV pairs are often mapped to server nodes via a hash function to balance storage or processing load over them. There is no standard for KV datastores; different KVs have differing constraints:

- Some offer ACID transactions
- Some allow very big objects in values, while others have hard limits
- Some constrain data-types of keys to be only string/numbers.

Some KV systems have built-in indexing of values:

- As KV pairs are added to DB, an index is compiled of unique values
- Produces a value->key LUT; index tells you all the keys for a given value.

Some KV systems auto-compress.

14.2 Document Databases

Document DBs manage more complex data structures than KV datastores. They are NOT stores of electronic documents.

Like KVs, DocDBs don't require you to define a common structure for all records in the data store. Unlike KVs, DocDBs allow you to query and filter collections of documents as you would rows in a RDB table.

In NoSQL, a 'document' is a structured object: it contains information not only about the data being stored, but also how that data is structured. A document in a DocDB is at root a set of KV pairs, but some of the values can themselves be sets of KV pairs. These documents within documents are called *embedded documents*

Most commonly, documents are described as JSON or XML files.

14.3 Columnar DBs

Column-family DBs, or columnar DBs, are used for the really big data. The technical term is VLDB, very large DB. Scaling for previously discussed DBs is possible, but are either hard, or extremely costly. Google's BigTable was the birth of 'column family' DBs.

‘Column family’ comes from the fact that groups of related columns that are frequently used together can be gathered into groups, called families. columns within a family are kept together on disk, to reduce access times.

The data-modeller defines the column *families* prior to implementing the DB, but devs can subsequently add columns to a family.

This means the designer of a DB specifies only a coarse grained structure. Column families are analogous to keys in KVDBs, with similar constraints wrt uniqueness of names in a namespace

Like DocDBs, ColFam DBs do not require column entries in all rows.

14.4 Graph DBs

Some forms of data are most naturally represented in a database as a set of nodes, edges joining nodes to represent relationships between them (many to many relationships that is difficult for RDBMs and SQL).

Specialised graph databases have been developed for efficient storage of such data.

Most graph algorithms are iterative, traverses the graph in some way, so MapReduce approach does not map onto graph processing very well.

Neo4J is a newer development that seems to be en route to becoming the de facto standard and outperform Apache’s Giraph.

15 Queues and Real-time Stream Processing

We have already talked about message queues, but here is a recap:

- Asynchronous
- Decoupling
- Resilience
- Redundancy
- Guarantees
- Scalable

However, message queues provide no transformation or processing.

Apache Kafka was originally used by LinkedIn, before being released as open source. Here are some of the features:

- Full independence of the producer and consumer
- Persistence
- Distributed and replicated
- Streams of records stored in categories called topics

- Guarantees:
 - Messages are appended in the order they are sent
 - Messages are consumed in the order they are stored
 - A replication factor N , meaning up to $N - 1$ server failures are tolerated without losing messages

Kafka has a log-based queue, and this is the basic persistence method used. Logs are sharded and replicated, ensuring topics are not limited to one node's capacity. It also uses a pub/sub strategy, but offers extra features, such as record persistence and replication, in a similar manner to many queues.

Kafka is critically dependent on Zookeeper. ZK is a distributed configuration and synchronisation service, where Kafka stores basic metadata about clusters, consumers, etc.

Kafka consumer groups can function as a queue. It also allows you to broadcast messages like pub-sub, but the Topic structure has both queue and pub/sub properties and can handle scale processing.

15.1 Real-time stream processing

Stream processing is the method by which data is processed from a continual stream of raw data, and can be used in things such as the Large Hadron Collider.

We want both high throughput and low latency, along with fault tolerance. RTSP is also 'straggler tolerant', where a straggler is an individual task who's execution time is taking much longer than expected.

A stream processor has a framework providing queueing or message passing (push/pull pipelines or pub/sub).

When considering RTSP, data needs to be considered:

- Data sources: How is the data stored? (binary/string)
- Size: How big is the data?
- Order

It is required that messages are delivered *exactly once* and messages are guaranteed to be delivered in the order they were sent. (always once and exactly once)

We can use Kafka for stream processing, too. There are three approaches to extend Kafka to handle this:

- Do it yourself
 - Looks simple
 - But is actually hard
- Fully fledged stream processing system
 - Apache Storm, Spark, Flink, Samza
- Lightweight stream-processing library
 - Kafka streams included in Kafka for ages, which is powerful and easy to use.

16 Spark

Complex jobs, interactive queries and online processing all need efficient primitives for data sharing, but in MapReduce, the only way to share data across jobs is stable storage (like the file system).

MapReduce would require a lot of I/O operations, and these calls could take up to 90% of the processing time. An alternative to this is *in-memory data sharing*. The basic idea is that we just write, well, in memory, rather than return to the disk every time. It can be between 10 or 100 times faster.

Spark and other **resilient distributed databases (RDDs)** is the technology that allows this.

RDD

An Rdd is:

- An immutable, partitioned, logical collection of records that can be stored in memory across the cluster
- May be partitioned by a key in each record
 - A key that partitions in a smart way improves performance
- May not be *materialized* until needed
- May be *cached* in memory if repeated use is anticipated (can also tell the system to cache data)

RDDs are created and manipulated through applying bulk transformations on disk or in memory (map, filter, join). Outputs to the driver node are generated by applying actions in parallel to RDDs (reduce, collect, count). RDDs are *lazy*; they don't materialize data until absolutely necessary.

Using an example of Log mining, we can load error messages from a log into memory and then interactively search for various patterns:

```
lines = spark.textFile("...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
messages.cache() // expected to use it again

messages.filter(_.contains("foo")).count
messages.filter(_.contains("bar")).count
```